



VIET NAM
NATIONAL
UNIVERSITY
HO CHI MINH CITY



PRESENTATION

AUTOMATIC DIFFERENTIATION

Presenters:

Nguyễn Quốc Bảo	2470076
Nguyễn Minh Chiến	2570179
Nguyễn Trọng Nhật	2570472
Nguyễn Đức Tài	2470321
Phan Thị Minh Thơ	2470112

Table of Contents

- 1 Introduction
- 2 Detaching Computation
- 3 Backpropagation
- 4 Gradients and Python Control Flow
- 5 Backward and forward comparison
- 6 Conclusion

Table of Contents

1 Introduction

2 Detaching Computation

3 Backpropagation

4 Gradients and Python Control Flow

5 Backward and forward comparison

6 Conclusion

Table of Contents

1 Introduction

2 **Detaching Computation**

3 Backpropagation

4 Gradients and Python Control Flow

5 Backward and forward comparison

6 Conclusion

Detaching Computation

Detaching treats a tensor as a constant during backpropagation by severing its link to the computational history

Suppose we have

$$z = x * y$$


and

$$y = x * x$$

In normal

$$\partial z / \partial x = 3x^2$$

```
x (input, requires_grad=True)
  |
  | (square operation)
  ↓
  y = x2
  |
  | (multiply by x)
  ↓
  z = x * y = x3
  |
  ↓
  ∂z/∂x = 3x2 (gradient flows
through entire graph)
```




Gradient: $\partial z / \partial x = 3x^2$

Detaching Computation

But we want to focus on the **direct** influence of x on z rather than the influence conveyed via y
=> we can create a new variable u that takes the same value as y but its ancestor is removed
=> $u.detach()$ will have $grad_fn = None$ (y have $grad_fn = MulBackward0$)

```
x (input, requires_grad=True)
  |
  | (square operation)
  ↓
y = x2
  |
  | .detach() ← BREAKS GRADIENT FLOW
  ↓
u = x2 NO grad_fn, treated as constant)
  |
  | (multiply by x)
  ↓
z = x * u
  |
  | ∂z/∂u = x (gradient ONLY through direct x,
  | NOT through u)
```




Gradient: $\partial z / \partial x = u = x^2$ (only direct path, treating u as constant)

Detaching Computation

But we want to focus on the **direct** influence of x on z rather than the influence conveyed via y
=> we can create a new variable u that takes the same value as y but its ancestor is removed
=> $u.detach()$ will have $grad_fn = None$ (y have $grad_fn = MulBackward0$)

```
x (input, requires_grad=True)
  |
  | (square operation)
  ↓
y = x2
  |
  | .detach() ← BREAKS GRADIENT FLOW
  ↓
u = x2 NO grad_fn, treated as constant)
  |
  | (multiply by x)
  ↓
z = x * u
  |
  | ∂z/∂u = x (gradient ONLY through direct x,
  | NOT through u)
```




Gradient: $\partial z / \partial x = u = x^2$ (only direct path, treating u as constant)

Table of Contents

- 1 Introduction
- 2 Detaching Computation
- 3 **Backpropagation**
- 4 Gradients and Python Control Flow
- 5 Backward and forward comparison
- 6 Conclusion

Backpropagation

Backpropagation is an algorithm used in neural networks to compute the gradients of the loss function with respect to the model's weights, by propagating the error backward through the network. These gradients are then used to update the weights using optimization methods like gradient descent.



Example

$$f(u, v) = (u + v)^2$$

$$u(a, b) = (a + b)^2, \quad v(a, b) = (a - b)^2$$

$$a(w, x, y, z) = (w + x + y + z)^2, \quad b(w, x, y, z) = (w + x - y - z)^2$$

Applying Chain Rule

$$\frac{\partial f}{\partial w} = \frac{\partial f}{\partial u} \frac{\partial u}{\partial w} + \frac{\partial f}{\partial v} \frac{\partial v}{\partial w}$$

$$\frac{\partial u}{\partial w} = \frac{\partial u}{\partial a} \frac{\partial a}{\partial w} + \frac{\partial u}{\partial b} \frac{\partial b}{\partial w}$$

$$\frac{\partial v}{\partial w} = \frac{\partial v}{\partial a} \frac{\partial a}{\partial w} + \frac{\partial v}{\partial b} \frac{\partial b}{\partial w}$$


Single differentiations

$$\frac{\partial f}{\partial u} = 2(u + v), \quad \frac{\partial f}{\partial v} = 2(u + v)$$

$$\frac{\partial u}{\partial a} = 2(a + b), \quad \frac{\partial u}{\partial b} = 2(a + b)$$

$$\frac{\partial v}{\partial a} = 2(a - b), \quad \frac{\partial v}{\partial b} = -2(a - b)$$

$$\frac{\partial a}{\partial w} = 2(w + x + y + z), \quad \frac{\partial b}{\partial w} = 2(w + x - y - z)$$



Backpropagation for non scalar variables

1. When `y` is a scalar

PyTorch can directly compute:

$$\frac{\partial y}{\partial x}$$

2. When `y` is a vector

The correct derivative is a **Jacobian matrix**:

$$J = \frac{\partial y}{\partial x}$$

This produces a **matrix**, not a vector.

3. Why PyTorch does not compute Jacobians automatically

Computing Jacobians is:

- Memory-intensive
- Slow

According to Pytorch document, if the tensor is non-scalar (i.e. its data has more than one element) and requires gradient, the function additionally requires specifying a gradient. It should be a tensor of matching type and shape, that represents the gradient of the differentiated function w.r.t. self.

In deep learning, we usually only need the **sum of gradients**, not the full Jacobian.

4. Why `.backward()` fails for vector outputs

PyTorch cannot guess how to reduce a vector to a scalar, so you must provide a gradient manually.



Example

We consider the input vector:

$$x = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}, \quad x \in \mathbb{R}^3$$

We define the vector-valued function:

$$y = f(x) = x^T x$$

That is:

$$y = \begin{bmatrix} x_1^2 \\ x_2^2 \\ x_3^2 \end{bmatrix}$$

The Jacobian of the function is:

$$J = \frac{\partial y}{\partial x} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} & \frac{\partial y_1}{\partial x_3} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} & \frac{\partial y_2}{\partial x_3} \\ \frac{\partial y_3}{\partial x_1} & \frac{\partial y_3}{\partial x_2} & \frac{\partial y_3}{\partial x_3} \end{bmatrix}$$

Since:

$$y_i = x_i^2$$

we have:

$$\frac{\partial y_i}{\partial x_j} = \begin{cases} 2x_i & i = j, \\ 0 & i \neq j. \end{cases}$$

Thus, the Jacobian is:

$$J = \begin{bmatrix} 2x_1 & 0 & 0 \\ 0 & 2x_2 & 0 \\ 0 & 0 & 2x_3 \end{bmatrix}$$

Example

We consider the input vector:

$$x = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}, \quad x \in \mathbb{R}^3$$

We define the vector-valued function:

$$y = f(x) = x^T x$$

That is:

```
... x: [1. 2. 3.]  
y: [1. 4. 9.]
```

---- PyTorch Behavior ----

Vector-Jacobian Product (v = [1,1,1]):
Result from backward(): [2. 4. 6.]

---- Manual Full Jacobian ----

	dy_1/dx_1	dy_2/dx_1	dy_3/dx_1
y_1	2.0	0.0	0.0
y_2	0.0	4.0	0.0
y_3	0.0	0.0	6.0

The Jacobian of the function is:

$$J = \frac{\partial y}{\partial x} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} & \frac{\partial y_1}{\partial x_3} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} & \frac{\partial y_2}{\partial x_3} \\ \frac{\partial y_3}{\partial x_1} & \frac{\partial y_3}{\partial x_2} & \frac{\partial y_3}{\partial x_3} \end{bmatrix}$$

Since:

$$y_i = x_i^2$$

we have:

$$\frac{\partial y_i}{\partial x_j} = \begin{cases} 2x_i & i = j, \\ 0 & i \neq j. \end{cases}$$

Thus, the Jacobian is:

$$J = \begin{bmatrix} 2x_1 & 0 & 0 \\ 0 & 2x_2 & 0 \\ 0 & 0 & 2x_3 \end{bmatrix}$$

Table of Contents

- 1 Introduction
- 2 Detaching Computation
- 3 Backpropagation
- 4 **Gradients and Python Control Flow**
- 5 Backward and forward comparison
- 6 Conclusion

1.1 Problem: Control Flow

Unlike Static Graph Frameworks, PyTorch allows differentiation through Python control structures like while, if, for.

Challenge: The function structure changes depending on input data. How to compute derivatives?

💡 **Question:** Is $f(a)$ a complex non-linear function?


PART 1: CONTROL FLOW

```
def f(a):
    """
    A function demonstrating dynamic control flow (while loop and if statement).
    """
    b = a * 2
    while b.norm() < 1000:
        b = b * 2

    if b.sum() > 0:
        c = b
    else:
        c = 100 * b
    return c
```

1.2 "Define-by-Run" Mechanism

PyTorch uses a **Dynamic Graph** mechanism. It doesn't pre-compile the while loop. Instead, it records the operations that *actually happen* during execution.



Loop Unrolling


For the Autograd engine, the loop above is simply a sequence of consecutive multiplications.

$$a \xrightarrow{\times 2} b_0 \xrightarrow{\times 2} b_1 \xrightarrow{\times 2} b_2 \cdots \xrightarrow{\times 2} b_n \cdots \rightarrow c$$

The system "forgets" the while logic and only remembers this operation sequence for Backpropagation.

It effectively calculates the gradient using the Chain Rule for this specific sequence of multiplications:

$$\frac{\partial c}{\partial a} = \frac{\partial c}{\partial b_n} \cdot \frac{\partial b_n}{\partial b_{n-1}} \cdots \frac{\partial b_0}{\partial a}$$



Gradient and Python Control Flow

1.3 Algebraic Analysis: Loops

Step 1: Analyzing the while loop

- Initialization: $b_0 = 2a$.
- Loop Logic: In each iteration, b is simply multiplied by 2.

After n iterations, the value of b will be:

$$b_n = 2^n \cdot (2a) = 2^{n+1} \cdot a$$

Where n is a positive integer that depends on the value of a (to satisfy the condition $\text{norm} < 1000$). However, for a specifically fixed a during a single forward pass, n is a constant.

Step 2: Analyzing the if statement

Finally, c will equal either b or $100b$.

In summary:

$$c = k \cdot a$$

Where k is an aggregate constant (it can be either 2^{n+1} or $100 \cdot 2^{n+1}$).

Step 3: Calculating the Derivative
The function is essentially a Linear Function with respect to the variable a :

$$f(a) = k \cdot a$$

The derivative of $f(a)$ with respect to a is:

$$\frac{df}{da} = k$$

On the other hand, from the function equation, we can extract k :

$$k = \frac{f(a)}{a}$$

Conclusion Therefore:

$$\frac{df}{da} = \frac{f(a)}{a}$$

This proves why $a.\text{grad} == d / a$ always returns True (within the limits of floating-point machine precision).

```
import torch

def f(a):
    """
    A function demonstrating dynamic control flow (while loop and if statement).
    """
    b = a * 2
    while b.norm() < 1000:
        b = b * 2

    if b.sum() > 0:
        c = b
    else:
        c = 100 * b
    return c

# We verify the derivative:
a = torch.randn(size=(1), requires_grad=True) # a is a random scalar/tensor
d = f(a)
d.backward()

# Check: Is a.grad equal to d / a?
print(f"Value of a: {a.item()}")
print(f"Value of f(a): {d.item()}")
print(f"Gradient calculated by python: {a.grad.item():10f}")
print(f"Gradient calculated by manual (f(a)/a): {d.item() / a.item():10f}")
print(f"a.grad == d/a: {a.grad == d / a}")

Value of a: -0.09854736179113388
Value of f(a): -161460.0
Gradient calculated by python: 1638400.000000
Gradient calculated by manual (f(a)/a): 1638400.024774
a.grad == d/a: True
```



1.4 Old Framework

```
import tensorflow.compat.v1 as tf
import numpy as np

# Tắt chế độ chạy ngay (Eager execution) để mô phỏng TF 1.x chuẩn
tf.disable_v2_behavior()

def f_static_graph(a_input):
    # 1. KHỞI TẠO BIẾN (PLACEHOLDER)
    # Trong Static Graph, ta phải tạo một "cái xô" rỗng để hứng dữ liệu sau này
    a = tf.placeholder(dtype=tf.float32, shape=(), name='a')

    # b = a * 2
    b_init = a * 2

    # 2. VÒNG LẶP WHILE (RẤT PHÚC TAP)
    # Phải định nghĩa hàm điều kiện (condition) và hàm thân vòng lặp (body) riêng biệt
    def condition(b):
        return tf.norm(b) < 1000

    def body(b):
        return b * 2

    # Dùng tf.while_loop để gắn node vòng lặp vào đồ thị
    b_loop = tf.while_loop(condition, body, [b_init])

    # 3. CÂU LỆNH IF/ELSE
    # Không được dùng "if b_loop.sum() > 0".
    # Phải dùng tf.cond và định nghĩa 2 hàm con cho 2 nhánh.
    def true_fn():
        return b_loop

    def false_fn():
        return 100 * b_loop

    c = tf.cond(tf.reduce_sum(b_loop) > 0, true_fn, false_fn)

    # 4. TÍNH GRADIENT
    # Phải khai báo graph tính đạo hàm ngay lúc xây dựng
    grads = tf.gradients(c, a)

    return a, c, grads
```

```
# --- PHẦN CHẠY CHƯƠNG TRÌNH (SESSION) ---
# Xây dựng đồ thị
a_ph, c_op, grad_op = f_static_graph(None)

# Tạo phiên làm việc (Session) để chạy đồ thị
with tf.Session() as sess:
    # Tạo dữ liệu giả lập (random input)
    input_val = np.random.randn()

    # CHẠY: Bom dữ liệu (feed_dict) vào placeholder để lấy kết quả
    c_val, grad_val = sess.run([c_op, grad_op], feed_dict={a_ph: input_val})

    print(f"Input a: {input_val}")
    print(f"Output f(a): {c_val}")
    print(f"Gradient: {grad_val[0]}")

    # Kiểm chứng lý thuyết: grad == f(a) / a
    print(f"Kiểm chứng (f(a)/a): {c_val / input_val}")

...
WARNING:tensorflow:From /usr/local/lib/python3.12/dist-packages/tensorflow/python/compat/v2_compat.py:98: disable_resource_variables is deprecated and will be removed in a future version.
Instructions for updating:
non-resource variables are not supported in the long term
Input a: 0.4885171644734144
Output f(a): 1000.483154296875
Gradient: 2048.0
Kiểm chứng (f(a)/a): 2048.0
```

1.5 Comparison

Feature	Legacy Frameworks (Static Graph - e.g., TF 1.x)	Modern Frameworks (Dynamic Graph - PyTorch)
Philosophy	Define-and-Run: You must define the entire graph structure first (build a fixed "factory") before feeding data into it.	Define-by-Run: Computation happens immediately as code is executed (on-the-fly). You calculate as you define.
Control Flow (Loops & Branching)	Complex: Requires specific APIs (e.g., <code>tf.while_loop</code> , <code>tf.cond</code>). Native Python logic (<code>if</code> , <code>while</code>) is ignored during graph compilation.	Simple: Uses native Python control flow (<code>if</code> , <code>while</code> , <code>for</code>) directly.
Debugging	Difficult: You cannot use <code>print()</code> or breakpoints inside the graph definition (body/condition functions) because the code is not actually running data yet.	Easy: You can use <code>print()</code> or standard debuggers inside loops to inspect intermediate values in real-time.
Gradient Calculation	Symbolic: You must explicitly define gradient nodes (<code>tf.gradients</code>) within the graph before starting the Session.	Imperative: You can call <code>.backward()</code> at any time after the forward pass is complete.
Coding Style	Declarative: The code can feel unintuitive or "foreign," distinct from standard Python programming.	Imperative: The code feels natural and "Pythonic," similar to writing standard algorithms.

Table of Contents

- 1 Introduction
- 2 Detaching Computation
- 3 Backpropagation
- 4 Gradients and Python Control Flow
- 5 **Backward and forward comparison**
- 6 Conclusion

Backward and forward comparison

Forward differentiation calculation ($x \rightarrow f$)

Example function:

$$f(x) = \log(x^2) \sin x + x^{-1}$$

Forward Differentiation ($x \rightarrow f$)


$$\begin{aligned}v'_1 &= \frac{dv_1}{dx} = \frac{d(x^2)}{dx} = 2x \\v'_2 &= \frac{dv_2}{dx} = \frac{dv_2}{dv_1} \cdot \frac{dv_1}{dx} = \frac{d(\log(v_1))}{dv_1} \cdot v'_1 = \frac{1}{v_1} \cdot 2x = \frac{2}{x} \\v'_3 &= \frac{dv_3}{dx} = \frac{d(\sin x)}{dx} = \cos x \\v'_4 &= (v_2 v_3)' = v'_2 v_3 + v_2 v'_3 = \frac{2}{x} \sin x + \log(x^2) \cos x \\v'_5 &= \frac{dv_5}{dx} = \frac{d(x^{-1})}{dx} = -\frac{1}{x^2}\end{aligned}$$

Derivative of (f):

$$\begin{aligned}f'(x) &= v'_4 + v'_5 \\f'(x) &= \frac{2}{x} \sin x + \log(x^2) \cos x - \frac{1}{x^2}\end{aligned}$$

Define the intermediate variables:

- $v_1 = x^2$
- $v_2 = \log(v_1)$
- $v_3 = \sin(x)$
- $v_4 = v_2 \cdot v_3$
- $v_5 = x^{-1}$
- $f = v_4 + v_5$



Backward and forward differentiation comparison

Backward differentiation calculation ($f \rightarrow x$)

Backward Differentiation ($f \rightarrow x$)

Start with:

$$\bar{f} = 1$$

$$\bar{v}_4 = \bar{f} \cdot \frac{\partial f}{\partial v_4} = 1 \cdot \frac{\partial(v_4 + v_5)}{\partial v_4} = 1 \cdot \left(\frac{\partial v_4}{\partial v_4} + \frac{\partial v_5}{\partial v_4} \right) = 1 \cdot (1 + 0) = 1$$

$$\bar{v}_5 = \bar{f} \cdot \frac{\partial f}{\partial v_5} = 1 \cdot \frac{\partial(v_4 + v_5)}{\partial v_5} = 1 \cdot \left(\frac{\partial v_4}{\partial v_5} + \frac{\partial v_5}{\partial v_5} \right) = 1 \cdot (0 + 1) = 1$$

$$\bar{v}_2 = \bar{v}_4 \cdot \frac{\partial v_4}{\partial v_2} = 1 \cdot \frac{\partial(v_2 \cdot v_3)}{\partial v_2} = 1 \cdot \left(\frac{\partial v_2}{\partial v_2} \cdot v_3 + v_2 \cdot \frac{\partial v_3}{\partial v_2} \right) = 1 \cdot (1 \cdot v_3 + v_2 \cdot 0) = v_3$$

$$\bar{v}_3 = \bar{v}_4 \cdot \frac{\partial v_4}{\partial v_3} = 1 \cdot \frac{\partial(v_2 \cdot v_3)}{\partial v_3} = 1 \cdot \left(v_2 \cdot \frac{\partial v_3}{\partial v_3} + v_3 \cdot \frac{\partial v_2}{\partial v_3} \right) = 1 \cdot (v_2 \cdot 1 + v_3 \cdot 0) = v_2$$

$$\bar{v}_1 = \bar{v}_2 \cdot \frac{\partial v_2}{\partial v_1} = v_3 \cdot \frac{\partial(\log v_1)}{\partial v_1} = v_3 \cdot \frac{1}{v_1} = \frac{v_3}{v_1}$$


$$\bar{x}_{v_5} = \bar{v}_5 \cdot \frac{\partial v_5}{\partial x} = 1 \cdot \left(-\frac{1}{x^2} \right) = -\frac{1}{x^2}$$

$$\bar{x}_{v_3} = \bar{v}_3 \cdot \frac{\partial v_3}{\partial x} = v_2 \cdot \cos x = \log(x^2) \cos x$$

$$\bar{x}_{v_1} = \bar{v}_1 \cdot \frac{\partial v_1}{\partial x} = \frac{v_3}{v_1} \cdot 2x = \frac{2}{x} \sin x$$

$$\bar{x} = \bar{x}_{v_1} + \bar{x}_{v_3} + \bar{x}_{v_5} = \frac{2}{x} \sin x + \log(x^2) \cos x - \frac{1}{x^2}$$

Final Gradient



Backward and forward differentiation comparison

Backward vs Forward Differentiation

Criteria	Forward differentiation	Backward differentiation
Direction	Inputs → Outputs	Outputs → Inputs
Computation Cost	\propto Inputs (n), good if $n \ll m$	\propto Outputs (m), good if $m \ll n$
Purpose	Compute gradients for inputs, suitable for functions with few inputs and many outputs	Compute gradients for many inputs, suitable for functions with many inputs and few outputs (e.g., loss function in ML)
Applications	Functions with few inputs and many outputs	Neural networks, machine learning, large-scale optimization
Implementation	Simple, no need to store intermediate values	More complex, requires storing intermediate values (tape)
Memory	Low	High (store intermediate values)

Table of Contents

1 Introduction

2 Detaching Computation

3 Backpropagation

4 Gradients and Python Control Flow

5 Backward and forward comparison

6 Conclusion

References

- THE END -

Thank you for your attention